

PathAI Engine - Assessment Report

Open University Learning Analytics Dataset (OULAD) | Engagement scoring, Week 6 risk prediction, and course recommendation

This report summarises a production-oriented PathAI prototype built to help universities understand engagement trajectories, identify students needing support by Week 6, and recommend suitable next modules. Each task produces an actionable product output rather than a standalone research result.

Task 1 - Behavioral Engagement Scoring

6 feature buckets	7 student archetypes	98.4% -> 19.0% risk gradient from lowest to highest score band
----------------------	-------------------------	---

The weekly 0-100 score uses activity volume, consistency, recency gaps, resource diversity, course-pace alignment, and assessment behavior. Archetypes translate those signals into recognizable patterns such as steady engagers, early dropouts, late recoverers, burst engagers, procrastinators, compliance engagers, and opportunistic engagers.

Task 2 - Disengagement Prediction And Staff Alerts

0.991 high-touch precision	0.870 ROC-AUC	0.793 top-60% F1	0.848 top-60% recall
-------------------------------	------------------	---------------------	-------------------------

The selected model is XGBoost, trained only on Week <= 6 features. Model selection and alerting were separated deliberately: XGBoost was chosen for ranking quality, while intervention thresholds were set by advisor capacity and alert-fatigue risk.

Task 3 - Course Recommendations

0.354 content hit@3	0.468 CF hit@3	7/7 catalog coverage	AAA, GGG, EEE cold-start modules
------------------------	-------------------	-------------------------	-------------------------------------

The recommender compares collaborative filtering against a shared feature-space content model. The content model matches student Week 6 profiles to course profiles using cosine similarity, blended with a Wilson lower-bound pass-rate prior for a conservative cold-start solution.

Overall result: the prototype is credible enough for a 90-day decision-support pilot. The strongest product insight is how model outputs are converted into workflows that advisors can act on without creating alert fatigue.

What This Ships Into Product

Student engagement view. A weekly 0-100 trajectory with an archetype label and the strongest behavioral signals behind the score. This lets administrators separate low grades from true disengagement.

Advisor risk queue. A ranked Week 6 intervention list split into high-touch support, light-touch nudges, and monitoring. This turns model risk into an operational workflow instead of a raw probability column.

Next-module guidance. Top-3 course recommendations with conservative cold-start defaults for students with no history. This gives students and advisors a concrete planning output for the next semester.

Primary caveat: the prototype is evaluated on historical OULAD outcomes, so the next 90 days should prioritize live advisor feedback, subgroup calibration, and recommendation constraints such as timetable availability.

Data Cleaning And Archetypes

The enrollment grain is id_student x code_module x code_presentation. Student ID alone is not unique in OULAD. Cleaning decisions preserve behavior signal while preventing target leakage.

Check	Issue Count	Policy
studentInfo key uniqueness	0	Use student-module-presentation as the enrolment key.
studentRegistration key uniqueness	0	Use student-module-presentation as the enrolment key.
Unregistered enrolments not labelled Withdrawn	9	Audit only; official final_result remains the supervised target.
Withdrawn enrolments without unregistration date	93	Do not use date_unregistration as a predictive feature because it can occur after Week 6.
Missing score rows	173	Keep missing scores as missing; median-impute only inside model preprocessing.
Invalid score rows outside 0-100	0	Invalid values would be excluded from score aggregates.
Assessment rows with missing due date	11	Exclude missing-date assessments from weekly due/submission features.
Negative VLE activity dates (pre-start engagement)	688988	Valid pre-presentation activity; date is days before module start. Used as proactivity signal.
VLE activity after module presentation length	0	Exclude from weekly engagement panel.
Missing VLE week_from/week_to metadata	10486	Do not use week_from/week_to as model features.

Candidate Feature	Week 6 Available?	Used?	Reason
Pre-start VLE activity (negative dates)	Yes	Yes	Observed before day 0; reflects eagerness to begin course materials.
Week 1-6 VLE clicks	Yes	Yes	Observed before the Week 6 prediction point.
Week 1-6 active days	Yes	Yes	Observed before the Week 6 prediction point.
Week 1-6 activity diversity	Yes	Yes	Observed before the Week 6 prediction point.
Early non-exam assessment submissions	Yes	Yes	Only due/submitted work by Week 6 is used.
Student demographics/context	Yes	Yes	Known at registration or course enrolment.
Engagement score	Yes	Yes	Computed from Week 1-6 behavior using train-derived score weights.
Final result	No	Target only	Used only to train/evaluate the supervised label.

Negative VLE dates were retained as pre-start proactivity rather than dropped. This provides useful early signal for both risk and recommendation tasks. In contrast, date_unregistration was audited but excluded from the Week 6 model because it can reveal future withdrawal.

Archetype	Behavioral Signature
Steady Engager	High mean engagement score, low score volatility, strong active-day consistency.
Early Dropout	Sharp negative early-to-late score delta, low late score, often withdrawn final result.
Late Recoverer	Large positive score delta from early weeks to late weeks.
Sporadic / Burst Engager	High score range, high max score, and several zero-click weeks.
Perfectionist Procrastinator	Successful/high-scoring enrolment with late submissions and lower punctuality.
Surface Level / Compliance Engager	High assessment completion but low activity diversity and low material-study frequency.
Opportunistic Engager	High material-click share and engagement spikes, but weaker active-week consistency.



Task 1: Engagement Score

Approach: construct weekly behavioral features, normalize by module-presentation-week peers, and aggregate with train-derived weights. Components with weak train signal are assigned zero weight.

Component	Success AUC	Final Weight
Peer-normalized click volume	0.774	15%
Module-material study frequency	0.741	15%
Activity diversity	0.762	15%
Recent activity	0.773	15%
Pre-start course material engagement	0.690	10%
Active-day consistency	0.778	10%
Assessment completion	0.718	10%
Low-weight assignment completion	0.684	10%
Three-week trend	0.481	0%
Assessment punctuality	0.469	0%

Feature	Correlation With Withdraw/Fail
engagement_score	-0.516
assessment_submitted_ratio	-0.486
zero_click_weeks_cum	+0.473
weekly_clicks_norm	-0.468
missed_assessments_cum	+0.453
inactive_last_14_days	+0.439

Key decision: prioritize recency, consistency, diversity, and assessment follow-through over raw click count. This keeps the score transparent and resistant to click-spam behavior.



Task 2: Week 6 Disengagement Model



Approach: leakage-safe Week 6 binary classification. Several models were compared on validation evidence, XGBoost was selected for ranking quality, and operating thresholds were then chosen from the precision-recall tradeoff based on advisor workload.

Model	Precision	Recall	F1	ROC-AUC
xgboost	0.738	0.842	0.787	0.863
xgboost_regularized	0.753	0.817	0.784	0.858
random_forest	0.754	0.816	0.783	0.857
logistic_regression	0.773	0.794	0.783	0.859

Cutoff	Precision	Recall	F1	Queue Share
High-touch 20%	0.991	0.377	0.546	20.1%
Top-60% light-touch	0.744	0.848	0.793	60.2%

Key decision: separate model quality from threshold policy. The model ranks risk; the product decides whether that risk becomes high-touch support, light-touch nudges, or monitoring. This avoids calling most students urgent while still preserving broad recall through the light-touch tier.

What an advisor receives: at the end of Week 6, each course presentation generates a ranked student list with student ID, course, risk score, tier, three plain-English reason codes, and a recommended action. High-touch students enter a live advisor queue and push notification; light-touch students enter a weekly digest and automated nudges; monitoring students remain visible without immediate action.



Task 3: Recommendation Engine

Approach: compare collaborative filtering to a shared feature-space content model. Student Week 6 vectors are matched to module profile vectors via cosine similarity, blended with a Wilson lower-bound pass prior.

Metric	Result	
Evaluation split	Train: 2013B/2013J/2014B, Holdout: 2014J	
Holdout students	1647	
Content hit@3	0.354	
Collaborative hit@3	0.468	
Content coverage	7/7	
Collaborative coverage	7/7	
Cold-start modules	AAA, GGG, EEE	
Actual Module	Content hit@3	CF hit@3
AAA	0.027	0.027
BBB	0.333	0.400
CCC	0.519	0.757
DDD	0.046	0.029
EEE	0.298	0.143
FFF	0.142	0.065
GGG	0.000	0.000

Key decision: replace raw pass-rate priors with Wilson lower-bound priors to reduce small-sample bias. This makes content recommendations more stable and defensible.

Gaps And 90-Day Plan

This prototype is intentionally simple and explainable, but not yet production-complete. The table below lists key gaps and concrete next actions.

Gap	Why It Matters	90-Day Action
Limited temporal depth	One-semester windows can miss drift	Move to rolling-semester retraining and drift checks
No advisor feedback loop	Model thresholds are not tied to intervention outcomes	Capture advisor action/outcome labels and recalibrate monthly
Recommendation label weakness	Next module may reflect timetable constraints	Add schedule constraints and student intent tags
Subgroup calibration not fully audited	Uneven risk quality can harm trust	Run subgroup calibration audits and apply targeted recalibration
Reproducibility hardening	Notebook and production states can diverge	Introduce versioned feature snapshots and run manifests
Phase	Delivery Focus	
Days 0-30	Deploy advisor dashboard in shadow mode; validate data freshness and triage quality	
Days 31-60	Tune intervention thresholds by module and staffing constraints	
Days 61-90	Retrain with feedback labels; publish impact and calibration report	

Why this is not best-in-class yet: it optimizes predictive performance on available labels, but does not yet optimize intervention utility, fairness constraints, and real scheduling feasibility end-to-end.